

Package ‘rENM.model’

April 26, 2026

Type Package

Title Modeling tools for the rENM framework

Version 0.1.0.9000

Author John L. Schnase

Maintainer John L. Schnase <rENM.Framework@gmail.com>

ORCID 0000-0001-7473-3977

Description The rENM framework is a modular sysem for reconstucting and analyzing long-term ecological niche dynamics using time-indexed environmental data and species occurrence records. Its packages support climate-based niche modeling, temporal analysis of climatic suitability, and reproducible workflows for studying ecological change.

License MIT + file LICENSE

Encoding UTF-8

RoxygenNote 7.3.3

URL <https://github.com/rENM-Framework/rENM.model>

BugReports <https://github.com/rENM-Framework/rENM.model/issues>

Imports rENM.core,

dismo,
doParallel,
dplyr,
foreach,
ggplot2,
maxnet,
raster,
scales,
sdm,
sf,
sp,
terra,
usdm

Contents

create_ensemble_model 2

create_range_map	4
create_timeseries	6
plot_suitability	8
reduce_covariance	10
save_suitability_plot	12
screen_by_convergence1	14
screen_by_convergence2	16
stage_all_variables	18
stage_occurrences	20
stage_screened_variables	21

Index	24
--------------	-----------

create_ensemble_model *Build an ensemble ENM and derived products for a species*

Description

Constructs an ensemble ecological niche model (ENM) for a given species (four-letter banding code) and a 5-year time slice using the sdm framework, producing ensemble predictions and derived products including range maps and model diagnostics.

Usage

```
create_ensemble_model(
  alpha_code,
  year,
  methods = c("maxnet", "rf", "brt", "glm", "mars"),
  reps = 3,
  bg = 2500,
  tp = 40,
  op = 2,
  ensemble_method = "unweighted",
  io = c("raster", "terra"),
  overwrite = TRUE,
  verbose = TRUE
)
```

Arguments

alpha_code	Character. Four-letter banding code, case-insensitive (validated as exactly four A-Z letters; stored upper-case).
year	Integer. A 5-year bin start in (1980, 2020) divisible by 5.
methods	Character. SDM algorithms to fit. Defaults to c("maxnet", "rf", "brt", "glm", "mars").
reps	Integer. Number of replicated subsampling runs. Default 3.
bg	Integer. Number of background points. Default 2500.
tp	Numeric. Test partition percent (0-100). Default 40.
op	Integer. Threshold optimization option passed to sdm::getEvaluation for binarization; 2 = max(se+sp). Default 2.

ensemble_method	Character. Ensemble combination method passed to <code>sdm::ensemble</code> in <code>setting = list(method = ...)</code> . Default "unweighted".
io	Character. Raster I/O backend for writing outputs; one of "raster" or "terra". Predictors are read via raster either way to maintain sdm compatibility. Default "raster".
overwrite	Logical. Overwrite existing output files. Default TRUE.
verbose	Logical. Emit timestamped console progress. Default TRUE.

Details

This function is part of the rENM framework's processing pipeline and operates within the project directory structure defined by `rENM_project_dir()`.

Pipeline context: Inputs/paths must follow the rENM run structure:

```
<rENM_project_dir()>/runs/<ALPHA>/TimeSeries/<YEAR>/occs/of-<YEAR>.csv
<rENM_project_dir()>/runs/<ALPHA>/TimeSeries/<YEAR>/vars/*.asc
<rENM_project_dir()>/runs/<ALPHA>/TimeSeries/<YEAR>/model/ (outputs written here)
```

Inputs: The occurrence CSV must include longitude, latitude.

Methods: Models are trained via `sdm::sdm()` with replicated subsampling (`reps`) and a test partition (`tp`, `percent`). Background points (`bg`) are drawn using `method = "gRandom"`.

The default modeling ensemble uses "maxnet" (a native R implementation of maximum entropy modeling) instead of "maxent" to avoid Java dependencies and improve reproducibility in CRAN-compliant environments.

Outputs:

- Ensemble prediction rasters (GeoTIFF and ASCII)
- Model statistics and variable importance summaries
- Binary range map derived via threshold optimization

Logging: A human-readable, append-only summary is written to: `<rENM_project_dir()>/runs/<alpha_code>/_log.txt` with key parameters, file outputs, and per-year status. Console messages include timestamps for traceability.

Value

Invisibly returns a List with structured elements:

- `data`: `sdmData` object containing training data
- `model`: `sdmModel` object representing fitted models
- `ensemble`: `RasterLayer` of predicted suitability
- `range`: `RasterLayer` of binary presence/absence
- `paths`: List of important written file paths
- `stats`: List of captured model statistics and variable importance outputs

Side effects include writing raster outputs, statistics files, variable importance summaries, plots, and appending a log entry to the run log file.

See Also

[sdm](#), [sdmData](#), [ensemble](#), [getEvaluation](#)

Examples

```
## Not run:
# Default knobs; raster I/O
result <- create_ensemble_model("CASP", 2000)

# Heavier replication, custom methods, terra I/O
result <- create_ensemble_model(
  alpha_code = "CASP", year = 2000,
  methods = c("maxnet", "glm", "brt", "rf", "mars"),
  reps = 5, bg = 5000, tp = 30, op = 2,
  ensemble_method = "unweighted",
  io = "terra", overwrite = TRUE, verbose = TRUE
)

## End(Not run)
```

create_range_map

Create and save a binary presence-absence range map for a species

Description

Build a ggplot2 map for a binary (0/1) presence-absence raster and save the plot image and raster outputs in multiple formats. Supports multiple raster input types and integrates species range overlays.

Usage

```
create_range_map(
  pa,
  alpha_code,
  year,
  state_shp = file.path(rENM_project_dir(),
    "data/shapefiles/tl_2012_us_state/tl_2012_us_state.shp"),
  title = "Range Map",
  layer = 1L,
  minor_div = 2L,
  width_px = 1000,
  height_px = 750,
  dpi = 72,
  title_size = 25,
  legend_size = 15,
  legend_title_size = 15,
  axis_title_size = 20,
  axis_text_size = 18
)
```

Arguments

pa	SpatRaster, Character, or Raster. A presence-absence raster.
alpha_code	Character. Code used to construct output paths and to look up GAP.RANGE via <code>get_species_info(alpha_code)</code> for locating the species range shapefile at: <code><rENM_project_dir()>/data/shapefiles/<gap_range>/<gap_range>.shp</code> . Must be a non-empty string (e.g., "casp").
year	Integer, Character. Year inserted into the output path and name.
state_shp	File path. Path to a states or administrative boundaries shapefile. Defaults to: <code><rENM_project_dir()>/data/shapefiles/tl_2012_us_state/tl_2012_us_state.shp</code> .
title	Character. Plot title. Default "Range Map".
layer	Integer. Layer index to use when pa has multiple layers. Default 1L.
minor_div	Integer. Number of minor grid divisions between major axis breaks. For example, 2L places two minor ticks between majors. Default 2L.
width_px, height_px	Numeric. Output PNG dimensions in pixels. Defaults 1000 and 750.
dpi	Numeric. Plot DPI used to convert pixels to inches for <code>ggsave()</code> . Default 72.
title_size, legend_size, legend_title_size, axis_title_size, axis_text_size	Numeric. Font sizes for plot elements.

Details

This function is part of the rENM framework's processing pipeline and operates within the project directory structure defined by `rENM_project_dir()`.

Pipeline context Constructs a presence-absence map and raster exports within the run-specific directory: `<rENM_project_dir()>/runs/<alpha_code>/TimeSeries/<year>/model/`

Inputs Presence (1) is rendered in orange; Absence (0) in light gray. NA cells are dropped from the plot. If pa has multiple layers, layer selects which to use.

Species range source (package helper) This function does not search for range shapefiles dynamically. It calls the internal helper `get_species_info(alpha_code)`, looks up GAP.RANGE in `<rENM_project_dir()>/data/_species.csv`, and loads the shapefile at: `<rENM_project_dir()>/data/shapefiles/<gap_range>.shp`

Methods Values in the raster are coerced to binary by rounding (`round()`) and then mapped to the factor levels "Presence" (1) and "Absence" (0) for plotting. The vector layers (states and species range) are reprojected to the raster CRS if necessary. Axis breaks are computed with `pretty()` plus configurable minor divisions.

Outputs The following files are written:

- Plot image: `<fn>.png`
- Raster (AAIGrid): `<fn>.asc`
- Raster (GeoTIFF): `<fn>.tif`

After writing .asc and .tif, common sidecar files (.prj, .aux.xml) are removed if present to keep the directory clean.

Value

Invisibly returns a List with:

- plot: ggplot2 object representing the rendered map
- img_path: Character path to the saved .png file
- asc_path: Character path to the saved .asc file
- tif_path: Character path to the saved .tif file

Side effects include writing raster and image files to disk and removing associated sidecar files when present.

See Also

[terra::rast()], [terra::writeRaster()], [sf::st_as_sf()], [ggplot2::ggplot()]

Examples

```
## Not run:
# Using a file path:
create_range_map(
  pa      = "rasters/casp_1980_pa.tif",
  alpha_code = "casp",
  year    = 1980
)

# Using a SpatRaster and a specific layer:
r <- terra::rast("rasters/casp_stack.tif")
create_range_map(r, "casp", 1990, layer = 2L,
  title = "CASP 1990 Range")

## End(Not run)
```

create_timeseries	<i>Build an ensemble rENM time series for a species</i>
-------------------	---

Description

Schedules and executes create_ensemble_model() for each 5-year time bin between 1980 and 2020 in parallel, producing a full rENM time series for a given species.

Usage

```
create_timeseries(alpha_code, project_dir = NULL)
```

Arguments

alpha_code	Character. Four-letter banding code, e.g., "CASP".
project_dir	Character, NULL. Base directory for outputs and logs. If NULL, the directory is resolved using rENM_project_dir().

Details

This function is part of the rENM.model framework's processing pipeline and operates within the project directory structure defined by `rENM_project_dir()`.

Inputs: A four-letter species code is required. The project directory is resolved either from the provided argument or via `rENM_project_dir()`.

Processing steps:

- Validates `alpha_code` and prepares the run directory
- Resolves the base project directory
- Spawns a PSOCK cluster with size: `min(number of years, detectCores())`
- Ensures each worker can call `create_ensemble_model()` within the rENM.model namespace
- Launches one job per 5-year bin using `parallel::parLapplyLB()` for load-balanced execution
- Collates per-year results including status, elapsed time, and output summaries
- Appends a formatted summary to the run log

Worker function availability: Workers in a PSOCK cluster are fresh R sessions. The function `create_ensemble_model()` is explicitly referenced from the rENM.model namespace and made available to workers via closure.

This supports both development (`devtools::load_all()`) and installed package workflows.

Fault tolerance: Each year is processed independently. Errors for individual years are captured and logged without aborting the full time series. The cluster is always stopped on exit.

Outputs: A processing summary is appended to:

```
file.path(project_dir, "runs", <alpha_code>, "_log.txt")
```

Value

Invisibly returns a named List keyed by year. Each element is:

- `ok`: Logical indicating successful completion
- `year`: Integer representing the 5-year bin
- `elapsed`: Numeric elapsed seconds for the task
- `result`: Output from `create_ensemble_model()` or NULL
- `notes`: Character containing "ok" or an error message

Attributes:

- `total_elapsed_sec`: Numeric total wall-clock time in seconds

Side effects:

- Executes parallel jobs using a PSOCK cluster
- Appends a processing summary to the run log
- Writes progress messages to the console

See Also

[`create_ensemble_model()`], **parallel**

Examples

```
## Not run:
  create_timeseries("CASP")
  create_timeseries("CASP", project_dir = "/my/project")

## End(Not run)
```

plot_suitability	<i>Plot a climatic suitability raster with state boundaries and species range</i>
------------------	---

Description

Build a ggplot2 map of continuous climatic suitability values from a raster and overlay state or administrative boundaries plus a species range polygon. If the raster has multiple layers, layer selects which to plot.

Usage

```
plot_suitability(
  en,
  alpha_code,
  state_shp = "~/rENM/data/shapefiles/tl_2012_us_state/tl_2012_us_state.shp",
  title = "Climatic Suitability Map",
  limits = c(0, 1),
  layer = 1L,
  minor_div = 2L
)
```

Arguments

en	SpatRaster, Character, Raster. A suitability raster. One of: a terra SpatRaster, a single-file path (character) to a raster readable by terra rast(), or a raster::Raster* object converted internally to SpatRaster.
alpha_code	Character. Species code used to look up GAP.RANGE via get_species_info(alpha_code) and, from that, the species range shapefile under <project_dir>/data/shapefiles/<gap_range>/<gap_range>.shp
state_shp	Character. Path to a states or administrative boundaries shapefile (.shp). Defaults to <project_dir>/data/shapefiles/tl_2012_us_state/ tl_2012_us_state.shp
title	Character. Plot title. Default "Climatic Suitability Map".
limits	Numeric. Length-2 vector giving the lower and upper limits of the color scale (default c(0, 1)).
layer	Integer. Band index to plot when en has multiple layers (default 1L).
minor_div	Integer. Number of minor grid divisions between major axis breaks (default 2L).

Details

This function is part of the rENM framework's processing pipeline and operates within the project directory structure defined by `rENM_project_dir()`.

Pipeline context Build a ggplot2 map of continuous climatic suitability values from a raster (e.g., 0-1) and overlay state or administrative boundaries plus a species range polygon. If the raster has multiple layers, `layer` selects which to plot.

The species range shapefile is looked up via an internal helper `get_species_info(alpha_code)`, which must return a component `GAP.RANGE`. The shapefile is then read from: `<project_dir>/data/shapefiles/<gap_range>.shp`

Inputs

- A suitability raster provided as a `SpatRaster`, file path, or raster object.
- Species code used to look up `GAP.RANGE` via `get_species_info(alpha_code)`.
- Path to a states or administrative boundaries shapefile.

Outputs

- A ggplot object representing the climatic suitability surface with state boundaries and species range overlays.

Methods

- The state and species-range vectors are reprojected to the raster CRS (if set) to ensure alignment.
- The color scale uses a smooth multi-hue palette via `grDevices::colorRampPalette()`.
- NA cell values render transparent.

Data requirements

- Species range shapefile located at: `<project_dir>/data/shapefiles/<gap_range>/<gap_range>.shp`
- Default state shapefile located at: `<project_dir>/data/shapefiles/tl_2012_us_state/tl_2012_us_state.shp`

Value

A ggplot object. Side effects:

- None. The function returns a plot object without writing files.

Examples

```
## Not run:  
# Example usage:  
plot_suitability(en, "CASP")  
  
## End(Not run)
```

reduce_covariance	<i>Reduce variable covariance</i>
-------------------	-----------------------------------

Description

Detects multicollinearity in the staged predictor rasters through an adaptive Variable Inflation Factor (VIF) routine, records results, and retains only non-correlated variables for further processing.

Usage

```
reduce_covariance(
  alpha_code,
  year = NULL,
  vif_threshold = 10,
  sample_schedule = c(10000, 20000, 50000),
  stability_patience = 1,
  corr_threshold = 0.95,
  exclude_vars = character(0),
  file_patterns = c("\\.tif$", "\\..grd$", "\\..img$", "\\..bil$"),
  overwrite_results = TRUE
)
```

Arguments

alpha_code	Character. Species alpha code (for example, "CASP").
year	Integer vector, or NULL. Years to process; if NULL (default) the sequence 1980, 1985, ..., 2020 is used.
vif_threshold	Numeric. VIF threshold passed to <code>usdm::vifstep()</code> (default 10).
sample_schedule	Integer vector. Adaptive sample sizes tried in order; the routine stops early once the kept set stabilizes. Default <code>c(10000, 20000, 50000)</code> .
stability_patience	Integer. Consecutive identical kept sets needed for convergence (default 1).
corr_threshold	Numeric. Absolute Pearson correlation threshold used when listing high-correlation pairs (default 0.95).
exclude_vars	Character vector. Predictor names that must always be retained. Case-sensitive match against sampled column names (default <code>character(0)</code>).
file_patterns	Character vector. Regular expressions selecting raster filenames to include. Matching stems trigger companion <code>.asc</code> moves. Default <code>c("\\.tif\$", "\\..grd\$", "\\..img\$", "\\..bil\$")</code> .
overwrite_results	Logical. Overwrite an existing results CSV (default TRUE).

Details

This function is part of the rENM framework's processing pipeline and operates within the project directory structure defined by `rENM_project_dir()`.

Pipeline context

- Input rasters: <project_dir>/runs/<alpha_code>/TimeSeries/ <year>/vars
- Log file: <project_dir>/runs/<alpha_code>/_log.txt

Inputs

- Predictor rasters matching file_patterns
- Optional exclusions via exclude_vars

Outputs

- Results CSV: _<alpha_code>-<year>-Correlation-Results.csv
- Directory of moved rasters (only if variables are dropped): _<alpha_code>-<year>-Correlated-Variables
- Log entry appended to _log.txt

Methods

- Adaptive sampling over sample_schedule sizes until the kept set is stable for stability_patience iterations
- Multicollinearity screening with usdm::vifstep() at threshold vif_threshold
- Reporting of high-correlation pairs where $|r| \geq corr_threshold$

Data requirements

- Raster layers must share identical geometry
- Layer names must correspond to predictor identifiers

Value

Invisible list keyed by year containing:

- kept – character vector of retained predictors
- dropped – character vector of removed predictors
- vif_table – data frame of VIF results for all variables
- cor_pairs – data frame of high-correlation pairs
- results_csv – file path to the written CSV
- moved_to – directory path for relocated rasters or NA when none were moved

Examples

```
## Not run:
reduce_covariance("CASP")
reduce_covariance("CASP", year = 2005,
                  exclude_vars = c("bio1", "bio12"))
reduce_covariance("CASP", year = 2000,
                  sample_schedule = c(8000, 16000, 32000, 64000))

## End(Not run)
```

save_suitability_plot *Save a ggplot suitability plot to file*

Description

Save a ggplot object to disk with configurable size, resolution, and typography. File format is inferred from the extension or defaults to PNG.

Usage

```
save_suitability_plot(
  p,
  fn,
  width = 1000,
  height = 750,
  res = 72,
  pointsize = 15,
  title_size = 28,
  subtitle_size = 20,
  legend_size = 15,
  legend_title_size = 15,
  axis_title_size = 20,
  axis_text_size = 18
)
```

Arguments

p	ggplot. A ggplot object to save.
fn	Character. Full output path and file name, with or without an extension. If no extension is provided, ".png" is appended.
width	Numeric. Image width in pixels. Default 1000.
height	Numeric. Image height in pixels. Default 750.
res	Numeric. Resolution (DPI). Default 72.
pointsize	Numeric. Base point size for text passed to ggsave(). Default 15.
title_size	Numeric. Font size for the plot title. Default 28.
subtitle_size	Numeric. Font size for the plot subtitle. Default 20.
legend_size	Numeric. Font size for legend text. Default 15.
legend_title_size	Numeric. Font size for the legend title. Default 15.
axis_title_size	Numeric. Font size for axis titles. Default 20.
axis_text_size	Numeric. Font size for axis tick labels. Default 18.

Details

This function is part of the rENM framework's processing pipeline and operates within the project directory structure defined by `rENM_project_dir()`.

Pipeline context This function saves ggplot outputs generated during the rENM workflow, resolving relative paths against the project directory defined by `rENM_project_dir()`.

Inputs

- A ggplot object.
- An output file path, with or without extension.

Outputs

- A saved image file written to disk.
- If no extension is provided, a PNG file is written.

Methods Width and height are specified in pixels and converted to inches using the resolution (DPI) for `ggplot2::ggsave()`.

Supported extensions:

- "png", "pdf", "tiff", "jpg", "jpeg", "svg", "eps"

Before saving, the plot theme is enhanced to apply specified font sizes to the title, subtitle, legend, and axis text elements.

Data requirements The output path is resolved relative to the project directory if a non-absolute path is provided.

Value

Invisibly returns the file path written.

- Character. File path to the saved plot.
- Side effect: Writes an image file to disk.

See Also

`[ggplot2::ggsave()]`, `[ggplot2::theme()]`, `[ggplot2::element_text()]`

Examples

```
## Not run:
p <- ggplot2::ggplot(mtcars, ggplot2::aes(mpg, wt)) +
  ggplot2::geom_point()
save_suitability_plot(p, "plots/mtcars_scatter.png")

# No extension provided -> defaults to PNG:
save_suitability_plot(p, "plots/mtcars_scatter")

## End(Not run)
```

screen_by_convergence1

Iterative bivariate ENM screening with adaptive, stability-aware convergence

Description

Repeatedly fits bivariate MaxEnt models, aggregates permutation importance (PI) scores, and stops when both the mean PI and membership of the adaptive top-k variable set stabilize. Designed for rapid, reproducible predictor down-selection across 5-year intervals.

Usage

```
screen_by_convergence1(
  alpha_code,
  year = NULL,
  max_iter = 50,
  batch_samples = NULL,
  w_stability = 0.5,
  background_n = 2500,
  seed = NULL,
  ncores = max(1, parallel::detectCores() - 1),
  fast = FALSE,
  set_threshold = 0.9,
  passes_required = NULL
)
```

Arguments

alpha_code	Character. Species or run code (e.g., "CASP").
year	Integer, Character, or NULL. Year to process; NULL runs all 5-year intervals 1980–2020.
max_iter	Integer. Maximum iterations safeguard; default 50.
batch_samples	Integer or NULL. Target appearances per variable per iteration; NULL triggers an adaptive fallback.
w_stability	Numeric. Instability weight in wz_score; default 0.5.
background_n	Integer. Background points per model; default 2500.
seed	Integer or NULL. RNG seed; NULL draws a random seed that is returned.
ncores	Integer. CPU cores for parallel fits; default <code>max(1, parallel::detectCores() - 1)</code> .
fast	Logical. Enable fast triage mode; default FALSE.
set_threshold	Numeric. Jaccard threshold for membership stability; default 0.90. Must lie in (0, 1).
passes_required	Integer or NULL. Consecutive passes needed for convergence; NULL defaults to 1 in fast mode, else 2.

Details

This function is part of the rENM framework's processing pipeline and operates within the project directory structure defined by `rENM_project_dir()`.

Pipeline context

- Reads occurrences from `runs/<alpha_code>/_occs/of-<year>.csv`
- Reads raster predictors from `runs/<alpha_code>/_vars/<year>/`
- Writes outputs to `runs/<alpha_code>/TimeSeries/<year>/vars/`

Inputs

- Occurrence table with longitude, latitude
- At least two GeoTIFF predictor layers

Methods Convergence requires both tests to pass:

- *Mean stability* – relative change in top-k mean PI below an adaptive threshold (10% if $p \leq 20$, 7.5% if $21 \leq p \leq 40$, else 5%).
- *Membership stability* – Jaccard overlap of consecutive top-k sets above `set_threshold`.

Adaptive top-k: $k = \min(12, \max(5, \text{round}(\sqrt{p})))$. Batch sampling defaults to $\max(1, \text{floor}(1.5 * \sqrt{p}))$ appearances per variable and is halved in fast mode. Parallel execution uses deterministic RNG streams.

Outputs

- Variable ranking CSV
- Convergence trace CSV
- Log entry appended to `_log.txt`

If `year = NULL`, all 5-year intervals 1980–2020 are processed.

Value

A list (invisible) with:

`importance_summary` Data frame of variable-wise PI statistics and rank diagnostics.

`convergence_trace` Data frame with per-iteration diagnostics.

`output_files` Named vector with paths to the two CSV files.

`seed_used` Integer seed actually used in this run.

Side effects: two CSV files and a log block are written to disk inside the project directory.

See Also

[maxent](#), [stack](#)

Examples

```
## Not run:
## Quick triage across all years
res <- screen_by_convergence1("CASP", fast = TRUE)

## Standard fidelity for a single year
res <- screen_by_convergence1("CASP", 1990, seed = 42)

## End(Not run)
```

```
screen_by_convergence2
```

Iterative bivariate ENM screening with adaptive, stability-aware convergence

Description

Repeatedly fits bivariate maxnet ecological niche models, aggregates permutation importance (PI) scores, and stops when both the mean PI and membership of the adaptive top-k variable set stabilize. Designed for rapid, reproducible predictor down-selection across 5-year intervals without the Java dependency required by `dismo::maxent()`.

Usage

```
screen_by_convergence2(
  alpha_code,
  year = NULL,
  max_iter = 50,
  batch_samples = NULL,
  w_stability = 0.5,
  background_n = 2500,
  seed = NULL,
  ncores = max(1, parallel::detectCores() - 1),
  fast = FALSE,
  set_threshold = 0.9,
  passes_required = NULL
)
```

Arguments

<code>alpha_code</code>	Character. Species or run code (e.g., "CASP").
<code>year</code>	Integer, Character, or NULL. Year to process; NULL runs all 5-year intervals 1980-2020.
<code>max_iter</code>	Integer. Maximum iterations safeguard; default 50.
<code>batch_samples</code>	Integer or NULL. Target appearances per variable per iteration; NULL triggers an adaptive fallback.
<code>w_stability</code>	Numeric. Instability weight in <code>wz_score</code> ; default 0.5.
<code>background_n</code>	Integer. Background points per model; default 2500.

seed	Integer or NULL. RNG seed; NULL draws a random seed that is returned.
ncores	Integer. CPU cores for parallel fits; default <code>max(1, parallel::detectCores() - 1)</code> .
fast	Logical. Enable fast triage mode; default FALSE.
set_threshold	Numeric. Jaccard threshold for membership stability; default 0.90. Must lie in $(0, 1)$.
passes_required	Integer or NULL. Consecutive passes needed for convergence; NULL defaults to 1 in fast mode, else 2.

Details

This function is part of the rENM framework's processing pipeline and operates within the project directory structure defined by `rENM_project_dir()`.

Pipeline context

- Reads occurrences from `runs/<alpha_code>/_occs/of-<year>.csv`
- Reads raster predictors from `runs/<alpha_code>/_vars/<year>/`
- Writes outputs to `runs/<alpha_code>/TimeSeries/<year>/vars/`

Inputs

- Occurrence table with longitude, latitude
- At least two GeoTIFF predictor layers

Methods The function repeatedly samples bivariate predictor pairs and fits maxnet models using occurrence and background environmental values extracted from the selected raster layers. Background points are drawn directly from non-NA raster cells using `raster::sampleRandom(xy = TRUE, na.rm = TRUE)`, eliminating the need for `dismo::randomPoints()`.

For each fitted pair, permutation importance is computed manually in a MaxEnt-like way: each predictor is permuted in turn on the model evaluation table, predictions are recomputed without re-fitting, the drop in training AUC is measured, and the resulting drops are normalized to percentages summing to 100 across the two predictors in that bivariate model. This produces a permutation-based importance metric analogous in spirit to Java MaxEnt's permutation importance, while remaining fully native to R.

Convergence requires both tests to pass:

- *Mean stability* - relative change in top-k mean PI below an adaptive threshold (10% if $p \leq 20$, 7.5% if $21 \leq p \leq 40$, else 5%).
- *Membership stability* - Jaccard overlap of consecutive top-k sets above `set_threshold`.

Adaptive top-k: $k = \min(12, \max(5, \text{round}(\sqrt{p})))$. Batch sampling defaults to `max(1, floor(1.5 * sqrt(p)))` appearances per variable and is halved in fast mode. Parallel execution uses deterministic RNG streams.

Outputs

- Variable ranking CSV
- Convergence trace CSV
- Log entry appended to `_log.txt`

If `year = NULL`, all 5-year intervals 1980-2020 are processed.

Value

A list (invisible) with:

importance_summary Data frame of variable-wise PI statistics and rank diagnostics.

convergence_trace Data frame with per-iteration diagnostics.

output_files Named vector with paths to the two CSV files.

seed_used Integer seed actually used in this run.

Side effects: two CSV files and a log block are written to disk inside the project directory.

See Also

[maxnet](#), [stack](#)

Examples

```
## Not run:
## Quick triage across all years
res <- screen_by_convergence2("CASP", fast = TRUE)

## Standard fidelity for a single year
res <- screen_by_convergence2("CASP", 1990, seed = 42)

## End(Not run)
```

stage_all_variables	<i>Stage environmental variables into TimeSeries bins</i>
---------------------	---

Description

Copies raster variables from run-level storage into TimeSeries bins, organizing predictors for downstream modeling while reporting progress and recording a processing summary in the run log.

Usage

```
stage_all_variables(
  alpha_code,
  year = seq(1980, 2020, by = 5),
  project_dir = NULL
)
```

Arguments

alpha_code	Character. Species alpha code (e.g., "CASP").
year	Integer. Vector of 5-year bins. Default: seq(1980, 2020, by = 5).
project_dir	Character, NULL. Optional project root directory. If NULL, resolved using rENM_project_dir().

Details

This function is part of the rENM framework's processing pipeline and operates within the project directory structure defined by `rENM_project_dir()`.

Inputs: Copies raster variables ('.tif', '.asc') from the run-level '_vars' directory into the corresponding TimeSeries directories:

```
<project_dir>/runs/<alpha_code>/_vars/<year>/*(tif|asc)
```

Outputs: Files are staged into:

```
<project_dir>/runs/<alpha_code>/TimeSeries/<year>/vars/*
```

Existing files are overwritten.

Methods:

- Recursively scans source directories for '.tif' and '.asc' files
- Copies files into the target directory (flat structure)
- Overwrites existing files
- Warns on missing source directories
- Warns on duplicate basenames (last file wins)

Logging: A processing summary is appended to:

```
<project_dir>/runs/<alpha_code>/_log.txt
```

including:

- Timestamp
- Alpha code
- Number of files transferred
- File list
- Total elapsed time
- Output location

Value

Data frame (invisibly returned) with columns:

- year
- src_dir
- dest_dir
- files_found
- files_copied
- duplicates

Side effects:

- Copies raster files into TimeSeries directories
- Overwrites existing files in target directories
- Appends a processing summary to the run log
- Writes progress messages to the console

Examples

```
## Not run:
stage_all_variables("CASP")
stage_all_variables("CASP", year = 2005)
stage_all_variables("CASP", year = c(1990, 2010, 2020))
stage_all_variables("CASP", project_dir = "/projects/rENM")

## End(Not run)
```

stage_occurrences	<i>Stage occurrence CSVs into time series bins</i>
-------------------	--

Description

Copies per-bin occurrence files into the TimeSeries structure for downstream modeling while ensuring consistency with preprocessing outputs and recording a detailed audit log.

Usage

```
stage_occurrences(alpha_code, year = seq(1980, 2020, 5), project_dir = NULL)
```

Arguments

alpha_code	Character. Species alpha code (e.g., "CASP").
year	Integer. Vector of 5-year bins. Default: seq(1980, 2020, 5).
project_dir	Character, NULL. Optional project root directory. If NULL, resolved using rENM_project_dir().

Details

This function is part of the rENM framework's processing pipeline and operates within the project directory structure defined by rENM_project_dir().

Inputs: Copies per-bin occurrence files (of-<year>.csv) from:

<project_dir>/runs/<alpha_code>/_occs/of-<year>.csv

Outputs: Files are staged into:

<project_dir>/runs/<alpha_code>/TimeSeries/<year>/occs/of-<year>.csv

Existing files are overwritten to ensure consistency with the latest preprocessing outputs.

Methods:

- Validates requested years (must be (1980, 2020), divisible by 5)
- Creates destination directories as needed
- Copies files using file.copy(..., overwrite = TRUE)
- Records per-year status: Copied, Missing, or Error
- Aggregates results into a structured log block

Logging: A standardized summary is appended to:

<project_dir>/runs/<alpha_code>/_log.txt

including:

- Timestamp and alpha code
- Per-year staging table
- Counts (total, copied, missing)
- Output paths (multi-line)
- Total elapsed time

Each run is separated by a 72-character divider for readability.

Value

List (invisibly returned) with elements:

- alpha_code: Character species alpha code
- years: Integer vector of processed years
- copied_paths: Character vector of copied file paths
- counts: List with elements total_years, valid_files, copied_files, missing_files
- elapsed_seconds: Numeric elapsed time in seconds
- log_file: Character path to the run log file

Side effects:

- Copies occurrence CSV files into TimeSeries directories
- Overwrites existing files in destination directories
- Appends a processing summary to the run log
- Writes audit output to the console

Examples

```
## Not run:
stage_occurrences("CASP")
stage_occurrences("CASP", year = 1990)
stage_occurrences("CASP", year = c(1980, 2000, 2015))

## End(Not run)
```

stage_screened_variables

Stage screened environmental variables into the time series

Description

Stage (copy) ranked raster predictor layers into per-year TimeSeries directories based on a master variable-ranking table. Supports several ranking metrics and an optional top-N filter.

Usage

```
stage_screened_variables(
  alpha_code,
  year = seq(1980, 2020, 5),
  top_n = 10,
  rank_method = "ratio"
)
```

Arguments

alpha_code	Character. Species alpha code (e.g., "CASP"). Used to resolve paths under the project runs directory.
year	Integer. Vector of years or single year. Defaults to seq(1980, 2020, 5). Each value must be a 5-year step within 1980-2020.
top_n	Integer or NULL. Number of top-ranked variables to stage per year. Default = 10; use NULL to copy all variables.
rank_method	Character. One of "ranksum", "medianPI", "ratio", "wz". Selects the ranking column used to order variables. Default = "ratio".

Details

This function is part of the rENM framework's processing pipeline and operates within the project directory structure defined by rENM_project_dir().

Inputs

- Master ranking CSV, one per year: <proj_root>/runs/<alpha_code>/TimeSeries/<year>/vars/_<alpha_code>-<year>-Variable-Ranking.csv
- Source rasters: <proj_root>/runs/<alpha_code>/_vars/<year>/

Outputs

- Ranked rasters copied to: <proj_root>/runs/<alpha_code>/TimeSeries/<year>/vars/
- Per-year status line and a 72-dash summary appended to: <proj_root>/runs/<alpha_code>/_log.txt

Methods

- Ranking columns: rank_ranksum, rank_medianPI, rank_ratio, rank_wz
- Tie-breakers use the associated score column (higher or lower depending on metric).

Data requirements

- Years must be multiples of five in the range {1980, 1985, ..., 2020}.
- Raster filenames must match the var column in the ranking CSV.

Value

Invisible named list with:

- alpha_code Character
- years Integer vector
- rank_method Character
- top_n Integer or NULL

- `copied_paths` Character vector
- `counts` List: `total_years`, `attempted_files`, `copied_files`, `missing_files`, `years_missing_ranking`
- `elapsed_seconds` Numeric
- `log_file` Character

Examples

```
## Not run:  
stage_screened_variables("CASP", top_n = 5, rank_method = "medianPI")  
  
## End(Not run)
```

Index

`create_ensemble_model`, [2](#)
`create_range_map`, [4](#)
`create_timeseries`, [6](#)

`ensemble`, [4](#)

`getEvaluation`, [4](#)

`maxent`, [15](#)
`maxnet`, [18](#)

`plot_suitability`, [8](#)

`reduce_covariance`, [10](#)

`save_suitability_plot`, [12](#)
`screen_by_convergence1`, [14](#)
`screen_by_convergence2`, [16](#)
`sdm`, [4](#)
`sdmData`, [4](#)
`stack`, [15](#), [18](#)
`stage_all_variables`, [18](#)
`stage_occurrences`, [20](#)
`stage_screened_variables`, [21](#)